



Chatbot con TensorFlow 2.0

Acecom IA



TensorFlow Blog

Artículo: Bryan M. Li, FOR.ai

Código: Abonia1



TensorFlow Blog

TensorFlow Blog

Search the Blog

Return to TensorFlow Home

All TensorFlow Core TensorFlow.js TensorFlow Lite TFX Community



DANFO.JS

Community · TensorFlow.js

Introducing Danfo.js, a Pandas-like Library in JavaScript

August 25, 2020 — A guest post by Rising Odegua, Independent Researcher; Stephen Oni, Data Science Nigeria



Japanese

Introducing TensorFlow Videos for a Global Audience: Japanese

August 24, 2020 — Posted by the TensorFlow Team

Tags

TensorFlow Core



TensorFlow.js



TensorFlow Lite



TFX



Community



TensorFlow.js

Introducing Semantic Reactor: Explore NLP in Google Sheets

August 21, 2020 — Posted by Dale Markowitz, Applied AI Engineer
Editor's note: An earlier version of this article was published on Dale's blog.



Procesamiento del dataset



Procesamiento del dataset

The Cornell Movie-Dialog

Descripción General:


- 220,579 conversaciones intercambiadas entre 10,292 personajes de películas.
- Incluye 9,035 personajes de 617 películas.
- En total 304,713 expresiones.



Procesamiento del dataset

The Cornell Movie-Dialog

Archivos:





Procesamiento del dataset

The Cornell Movie-Dialog

1. movie_conversations.txt:

```
u0 +++$+++ u2 +++$+++ m0 +++$+++ ['L194', 'L195', 'L196', 'L197']  
u0 +++$+++ u2 +++$+++ m0 +++$+++ ['L198', 'L199']  
u0 +++$+++ u2 +++$+++ m0 +++$+++ ['L200', 'L201', 'L202', 'L203']  
u0 +++$+++ u2 +++$+++ m0 +++$+++ ['L204', 'L205', 'L206']  
u0 +++$+++ u2 +++$+++ m0 +++$+++ ['L207', 'L208']
```



Procesamiento del dataset

The Cornell Movie-Dialog

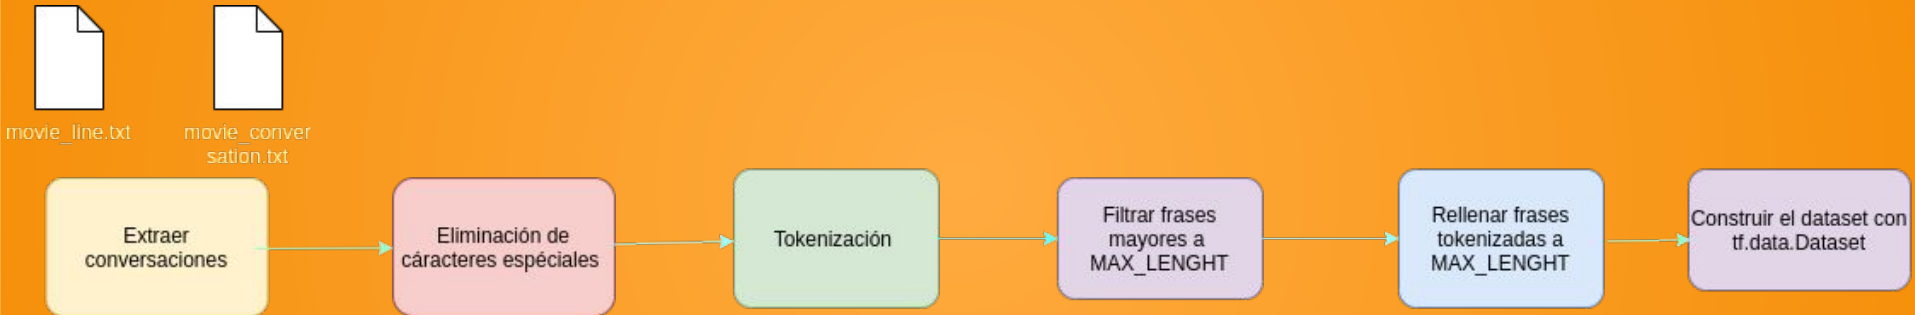
1. movie_lines.txt:

```
L901 +++$+++ u5 +++$+++ m0 +++$+++ KAT +++$+++ He said everyone was doing it. So I did it.  
L900 +++$+++ u0 +++$+++ m0 +++$+++ BIANCA +++$+++ As in..  
L899 +++$+++ u5 +++$+++ m0 +++$+++ KAT +++$+++ Now I do. Back then, was a different story.  
L898 +++$+++ u0 +++$+++ m0 +++$+++ BIANCA +++$+++ But you hate Joey  
L897 +++$+++ u5 +++$+++ m0 +++$+++ KAT +++$+++ He was, like, a total babe
```




Procesamiento del dataset

Pipeline





Procesamiento del dataset

- Path_line: ruta de movie_line.txt
- Path_movie_conversation: ruta de movie_conversation.txt
- Max_Samples: Max cantidad de ejemplos de entrenamiento(50000)

```
def preprocess_sentences(frase):
    frase = frase.lower().strip()
    frase = re.sub(r"([?.,!])", r"\1 ", frase)
    frase = re.sub(r'[" "]+' , " ", frase)
    frase = re.sub(r"^[a-zA-Z?.,!]+", " ", frase)
    frase = frase.strip()
    return frase

def load_sentence():
    id2line = {}
    with open(path_lines, errors="ignore") as file:
        lines = file.readlines()
        for line in lines:
            parts = line.replace("\n", "").split(" +++$+++ ")
            id2line[parts[0]] = parts[4]

    inputs, outputs = [], []
    with open(path_movie_conversations, "r") as file:
        lines = file.readlines()
        for line in lines:
            parts = line.replace("\n", "").split(" +++$+++ ")
            conversation = [line[1:-1] for line in parts[3][1:-1].split(", ")]
            for i in range(len(conversation) - 1):
                inputs.append(preprocess_sentences(id2line[conversation[i]]))
                outputs.append(preprocess_sentences(id2line[conversation[i + 1]]))
            if len(inputs) >= MAX_SAMPLES:
                return inputs, outputs
    return inputs, outputs
```



Procesamiento del dataset

Which means unknown word pieces will be encoded one character at a time. It's best understood through an example. Let's suppose you build a `SubwordTextEncoder` using a very large corpus of English text such that most of the common words are in vocabulary.

```
vocab_size = 10000
tokenizer = tfds.features.text.SubwordTextEncoder.build_from_corpus(
    corpus_sentences, vocab_size)
```

Let's say you try to tokenize the following sentence.

```
tokenizer.encode("good badwords badxyz")
```

It will be tokenized as:

1. good
2. bad
3. words
4. bad
5. x
6. y
7. z

As you can see, since the word piece "xyz" is not in vocabulary it is tokenized as characters.

```
# Build tokenizer using tfds for both questions and answers
tokenizer = tfds.features.text.SubwordTextEncoder.build_from_corpus(
    inputs + outputs, target_vocab_size=2*13)

START_TOKEN, END_TOKEN = [tokenizer.vocab_size], [tokenizer.vocab_size + 1]
VOCAB_SIZE = tokenizer.vocab_size + 2

print(tokenizer.encode("hello there!!"))
print(tokenizer.decode(tokenizer.encode("hello there!!")))
print(tokenizer.vocab_size)

... [2276, 126, 8110, 8110]
hello there!!
8333

print(START_TOKEN, END_TOKEN)

... [8333] [8334]
```



Procesamiento del dataset

- Max_LENGTH: Max longitud de la frase.(40)

```
def tokenize_and_filter(input, output, MAX_LENGTH):
    tokenize_inputs, tokienize_outputs = [], []
    for(sentence1, sentence2) in zip(input, output):
        sentence1 = START_TOKEN + tokenizer.encode(sentence1) + END_TOKEN
        sentence2 = START_TOKEN + tokenizer.encode(sentence2) + END_TOKEN
        if len(sentence1) <= MAX_LENGTH and len(sentence2) <= MAX_LENGTH:
            tokenize_inputs.append(sentence1)
            tokienize_outputs.append(sentence2)
    tokenize_inputs = tf.keras.preprocessing.sequence.pad_sequences(
        tokenize_inputs, maxlen= MAX_LENGTH, padding= 'post')
    tokienize_outputs = tf.keras.preprocessing.sequence.pad_sequences(
        tokienize_outputs, maxlen= MAX_LENGTH, padding= 'post')
    return tokenize_inputs, tokienize_outputs
```



Procesamiento del dataset

```
BATCH_SIZE = 64
BUFFER_SIZE = 20000

dataset = tf.data.Dataset.from_tensor_slices((
    {
        'inputs': inputs,
        'dec_inputs': outputs[:, :-1]
    },
    {
        'outputs': outputs[:, 1:]
    },
))

dataset = dataset.cache()
dataset = dataset.shuffle(BUFFER_SIZE)
dataset = dataset.batch(BATCH_SIZE)
dataset = dataset.prefetch(tf.data.experimental.AUTOTUNE)
```

Python

```
dataset
```

Python

```
<PrefetchDataset shapes: ({inputs: (None, 40), dec_inputs: (None, 39)}, {outputs: (None, 39)}), types: ({inputs: tf.int32,
dec_inputs: tf.int32}, {outputs: tf.int32})>
```



Transformers



Arquitectura Transformer

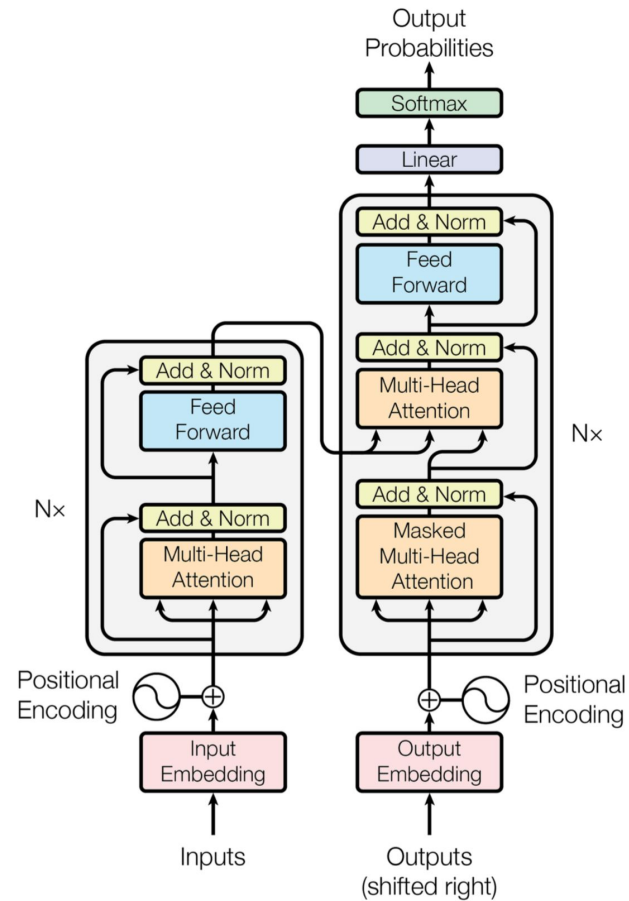


Figure 1: The Transformer - model architecture.




Implementación Multi-Head Attention



Implementación de Atención

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

scaled





Implementación de Atención

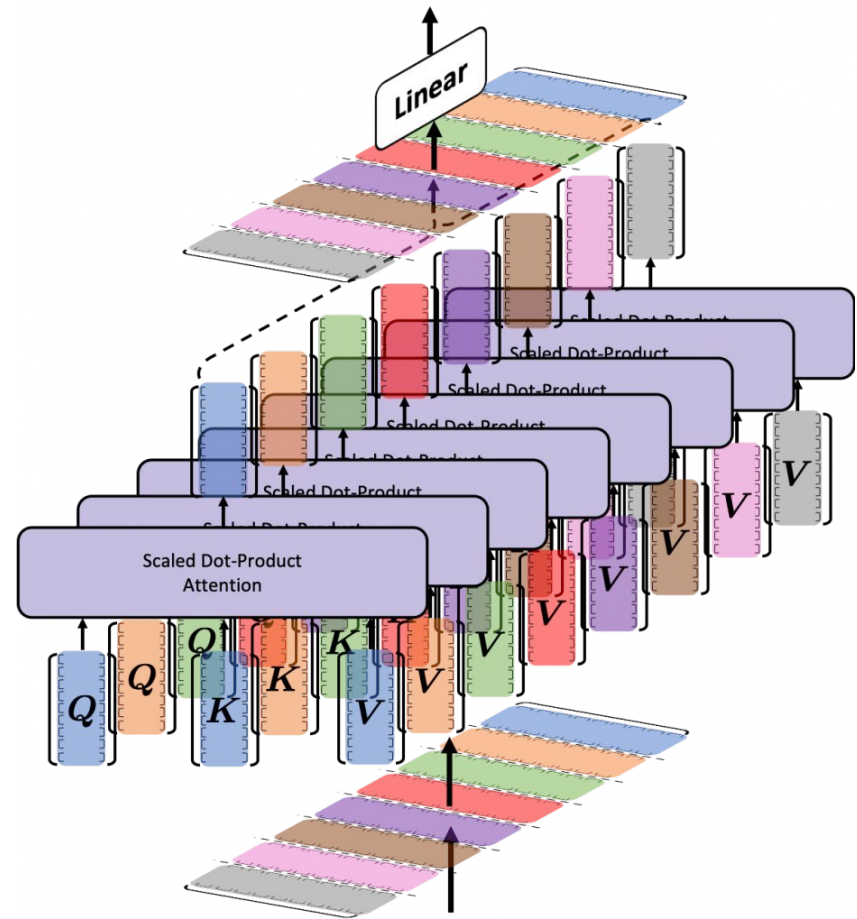
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

The term $\frac{QK^T}{\sqrt{d_k}}$ is annotated with a red box around the denominator $\sqrt{d_k}$ and a red arrow pointing to it from the word "scaled" written in red text.

```
def scaled_dot_product_attention(query, key, value, mask):  
    matmul_qk = tf.matmul(query, key, transpose_b=True)  
  
    depth = tf.cast(tf.shape(key)[-1], tf.float32)  
    logits = matmul_qk / tf.math.sqrt(depth)  
  
    # add the mask zero out padding tokens.  
    if mask is not None:  
        logits += (mask * -1e9)  
  
    attention_weights = tf.nn.softmax(logits, axis=-1)  
  
    return tf.matmul(attention_weights, value)
```



Implementación de MultiHead Attention





Implementación de Multi Head Attention

```
class MultiHeadAttention(tf.keras.layers.Layer):  
  
    def __init__(self, d_model, num_heads, name="multi_head_attention"):  
        super(MultiHeadAttention, self).__init__(name=name)  
        self.num_heads = num_heads  
        self.d_model = d_model  
  
        assert d_model % self.num_heads == 0  
  
        self.depth = d_model // self.num_heads  
  
        self.query_dense = tf.keras.layers.Dense(units=d_model)  
        self.key_dense = tf.keras.layers.Dense(units=d_model)  
        self.value_dense = tf.keras.layers.Dense(units=d_model)  
  
        self.dense = tf.keras.layers.Dense(units=d_model)  
  
    def split_heads(self, inputs, batch_size):  
        inputs = tf.reshape(  
            inputs, shape=(batch_size, -1, self.num_heads, self.depth))  
        return tf.transpose(inputs, perm=[0, 2, 1, 3])
```



Implementación de Multi Head Attention

```
def call(self, inputs):
    query, key, value, mask = inputs['query'], inputs['key'], inputs[
        'value'], inputs['mask']
    batch_size = tf.shape(query)[0]

    # linear layers
    query = self.query_dense(query)
    key = self.key_dense(key)
    value = self.value_dense(value)

    # split heads
    query = self.split_heads(query, batch_size)
    key = self.split_heads(key, batch_size)
    value = self.split_heads(value, batch_size)

    scaled_attention = scaled_dot_product_attention(query, key, value, mask)

    scaled_attention = tf.transpose(scaled_attention, perm=[0, 2, 1, 3])

    concat_attention = tf.reshape(scaled_attention,
                                   (batch_size, -1, self.d_model))

    outputs = self.dense(concat_attention)

    return outputs
```



Enmascaramiento

create_paddind_mask

```
def create_padding_mask(x):  
    mask = tf.cast(tf.math.equal(x, 0), tf.float32)  
    # (batch_size, 1, 1, sequence length)  
    return mask[:, tf.newaxis, tf.newaxis, :]
```

[]

```
def create_look_ahead_mask(x):  
    seq_len = tf.shape(x)[1]  
    look_ahead_mask = 1 - tf.linalg.band_part(tf.ones((seq_len, seq_len)), -1, 0)  
    padding_mask = create_padding_mask(x)  
    return tf.maximum(look_ahead_mask, padding_mask)
```

[]

```
: print(create_padding_mask(tf.constant([[1, 2, 0, 3, 0], [0, 0, 0, 4, 5]])))
```

```
tf.Tensor(  
[[[0. 0. 1. 0. 1.]]]
```

```
[[[1. 1. 1. 0. 0.]]]], shape=(2, 1, 1, 5), dtype=float32)
```



Enmascaramiento

create_look_ahead_mask

```
def create_padding_mask(x):  
    mask = tf.cast(tf.math.equal(x, 0), tf.float32)  
    # (batch_size, 1, 1, sequence length)  
    return mask[:, tf.newaxis, tf.newaxis, :]
```

[]

```
def create_look_ahead_mask(x):  
    seq_len = tf.shape(x)[1]  
    look_ahead_mask = 1 - tf.linalg.band_part(tf.ones((seq_len, seq_len)), -1, 0)  
    padding_mask = create_padding_mask(x)  
    return tf.maximum(look_ahead_mask, padding_mask)
```

[]


```
: print(create_look_ahead_mask(tf.constant([[1, 2, 0, 4, 5]])))
```

```
tf.Tensor(  
[[[0. 1. 1. 1. 1.]  
  [0. 0. 1. 1. 1.]  
  [0. 0. 1. 1. 1.]  
  [0. 0. 1. 0. 1.]  
  [0. 0. 1. 0. 0.] ]], shape=(1, 1, 5, 5), dtype=float32)
```



Positional encoding

- d es la dimensión de los word embedding
- pos es la posición de la palabra.
- i se refiere a cada una de las diferentes dimensiones de los positional encoding



$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{model}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{model}})$$

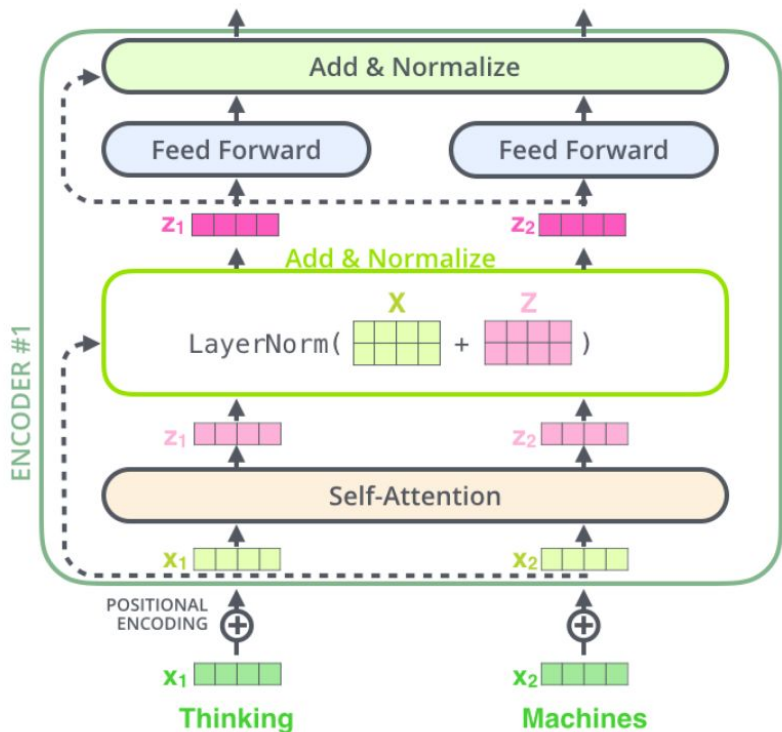


Positional encoding

```
class PositionalEncoding(tf.keras.layers.Layer):  
  
    def __init__(self, position, d_model):  
        super(PositionalEncoding, self).__init__()  
        self.pos_encoding = self.positional_encoding(position, d_model)  
  
    def get_angles(self, position, i, d_model):  
        angles = 1 / tf.pow(10000, (2 * (i // 2)) / tf.cast(d_model, tf.float32))  
        return position * angles  
  
    def positional_encoding(self, position, d_model):  
        angle_rads = self.get_angles(  
            position=tf.range(position, dtype=tf.float32)[:], tf.newaxis],  
            i=tf.range(d_model, dtype=tf.float32)[tf.newaxis, :],  
            d_model=d_model)  
        # apply sin to even index in the array  
        sines = tf.math.sin(angle_rads[:, 0::2])  
        # apply cos to odd index in the array  
        cosines = tf.math.cos(angle_rads[:, 1::2])  
  
        pos_encoding = tf.concat([sines, cosines], axis=-1)  
        pos_encoding = pos_encoding[tf.newaxis, ...]  
        return tf.cast(pos_encoding, tf.float32)  
  
    def call(self, inputs):  
        return inputs + self.pos_encoding[:, :tf.shape(inputs)[1], :]
```



Encoding layer



```
def encoder_layer(units, d_model, num_heads, dropout, name="encoder_layer"):
    inputs = tf.keras.Input(shape=(None, d_model), name="inputs")
    padding_mask = tf.keras.Input(shape=(1, 1, None), name="padding_mask")

    attention = MultiHeadAttention(
        d_model, num_heads, name="attention")({
            'query': inputs,
            'key': inputs,
            'value': inputs,
            'mask': padding_mask
        })

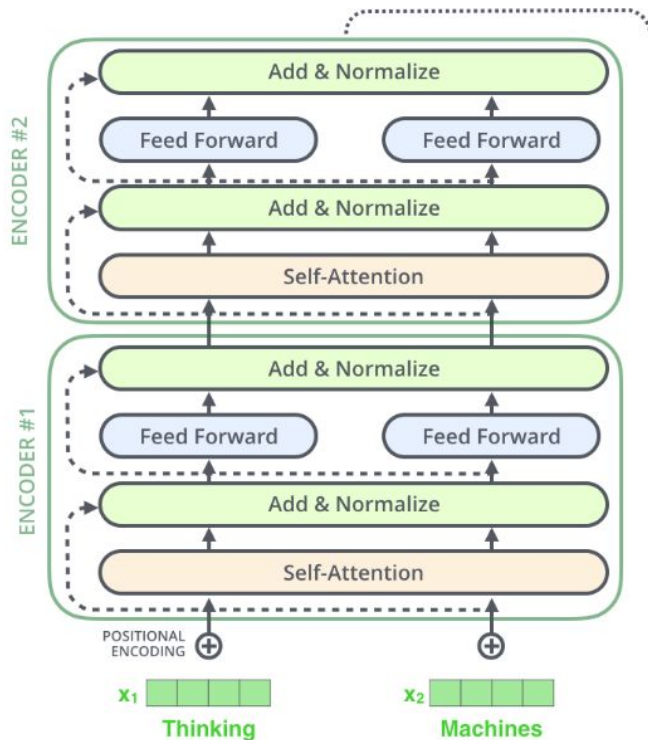
    attention = tf.keras.layers.Dropout(rate=dropout)(attention)
    attention = tf.keras.layers.LayerNormalization(
        epsilon=1e-6)(inputs + attention)

    outputs = tf.keras.layers.Dense(units=units, activation='relu')(attention)
    outputs = tf.keras.layers.Dense(units=d_model)(outputs)
    outputs = tf.keras.layers.Dropout(rate=dropout)(outputs)
    outputs = tf.keras.layers.LayerNormalization(
        epsilon=1e-6)(attention + outputs)

    return tf.keras.Model(
        inputs=[inputs, padding_mask], outputs=outputs, name=name)
```



Encoder



```
def encoder(vocab_size,
            num_layers,
            units,
            d_model,
            num_heads,
            dropout,
            name="encoder"):
    inputs = tf.keras.Input(shape=(None,), name="inputs")
    padding_mask = tf.keras.Input(shape=(1, 1, None), name="padding_mask")

    embeddings = tf.keras.layers.Embedding(vocab_size, d_model)(inputs)
    embeddings *= tf.math.sqrt(tf.cast(d_model, tf.float32))
    embeddings = PositionalEncoding(vocab_size, d_model)(embeddings)


    outputs = tf.keras.layers.Dropout(rate=dropout)(embeddings)

    for i in range(num_layers):
        outputs = encoder_layer(
            units=units,
            d_model=d_model,
            num_heads=num_heads,
            dropout=dropout,
            name="encoder_layer_{}".format(i),
        )([outputs, padding_mask])

    return tf.keras.Model(
        inputs=[inputs, padding_mask], outputs=outputs, name=name)
```



Decoding layer



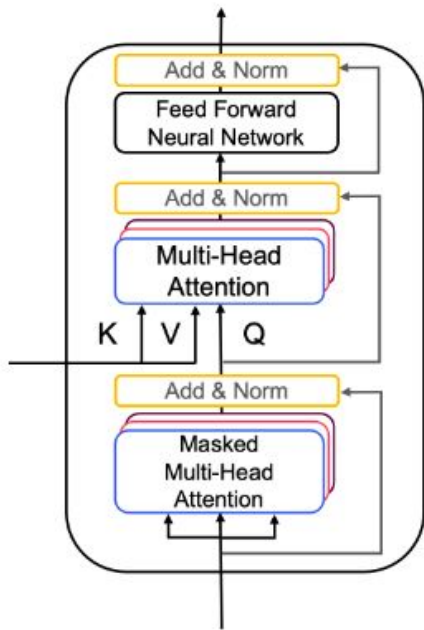
```
def decoder_layer(units, d_model, num_heads, dropout, name="decoder_layer"):
    inputs = tf.keras.Input(shape=(None, d_model), name="inputs")
    enc_outputs = tf.keras.Input(shape=(None, d_model), name="encoder_outputs")
    look_ahead_mask = tf.keras.Input(
        shape=(1, None, None), name="look_ahead_mask")
    padding_mask = tf.keras.Input(shape=(1, 1, None), name='padding_mask')

    attention1 = MultiHeadAttention(
        d_model, num_heads, name="attention_1")(inputs={
            'query': inputs,
            'key': inputs,
            'value': inputs,
            'mask': look_ahead_mask
        })
    attention1 = tf.keras.layers.LayerNormalization(
        epsilon=1e-6)(attention1 + inputs)

    attention2 = MultiHeadAttention(
        d_model, num_heads, name="attention_2")(inputs={
            'query': attention1,
            'key': enc_outputs,
            'value': enc_outputs,
            'mask': padding_mask
        })
```



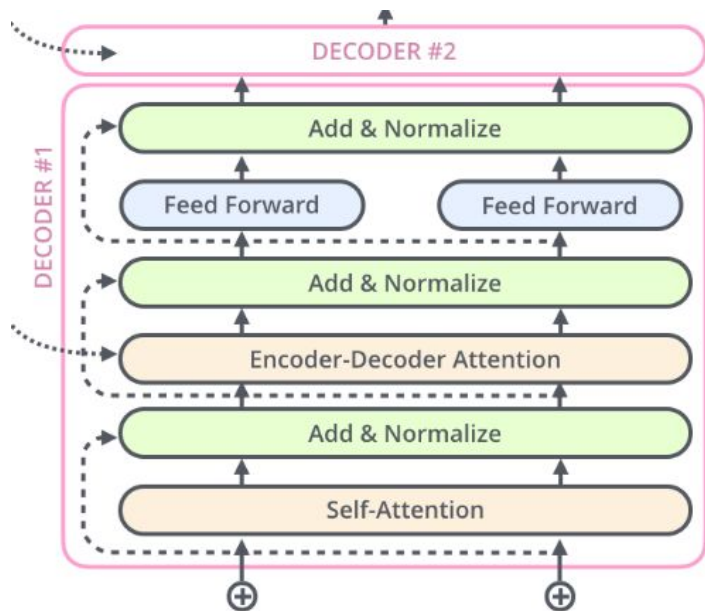
Decoding layer



```
attention1 = tf.keras.layers.LayerNormalization(  
    epsilon=1e-6)(attention1 + inputs)  
  
attention2 = MultiHeadAttention(  
    d_model, num_heads, name="attention_2")(inputs={  
        'query': attention1,  
        'key': enc_outputs,  
        'value': enc_outputs,  
        'mask': padding_mask  
    })  
attention2 = tf.keras.layers.Dropout(rate=dropout)(attention2)  
attention2 = tf.keras.layers.LayerNormalization(  
    epsilon=1e-6)(attention2 + attention1)  
  
outputs = tf.keras.layers.Dense(units=units, activation='relu')(attention2)  
outputs = tf.keras.layers.Dense(units=d_model)(outputs)  
outputs = tf.keras.layers.Dropout(rate=dropout)(outputs)  
outputs = tf.keras.layers.LayerNormalization(  
    epsilon=1e-6)(outputs + attention2)  
  
return tf.keras.Model(  
    inputs=[inputs, enc_outputs, look_ahead_mask, padding_mask],  
    outputs=outputs,  
    name=name)
```



Decoder




```
def decoder(vocab_size,
            num_layers,
            units,
            d_model,
            num_heads,
            dropout,
            name='decoder'):
    inputs = tf.keras.Input(shape=(None,), name='inputs')
    enc_outputs = tf.keras.Input(shape=(None, d_model), name='encoder_outputs')
    look_ahead_mask = tf.keras.Input(
        shape=(1, None, None), name='look_ahead_mask')
    padding_mask = tf.keras.Input(shape=(1, 1, None), name='padding_mask')

    embeddings = tf.keras.layers.Embedding(vocab_size, d_model)(inputs)
    embeddings *= tf.math.sqrt(tf.cast(d_model, tf.float32))
    embeddings = PositionalEncoding(vocab_size, d_model)(embeddings)

    outputs = tf.keras.layers.Dropout(rate=dropout)(embeddings)
```




Decoder



```
for i in range(num_layers):
    outputs = decoder_layer(
        units=units,
        d_model=d_model,
        num_heads=num_heads,
        dropout=dropout,
        name='decoder_layer_{}'.format(i),
    )(inputs=[outputs, enc_outputs, look_ahead_mask, padding_mask])

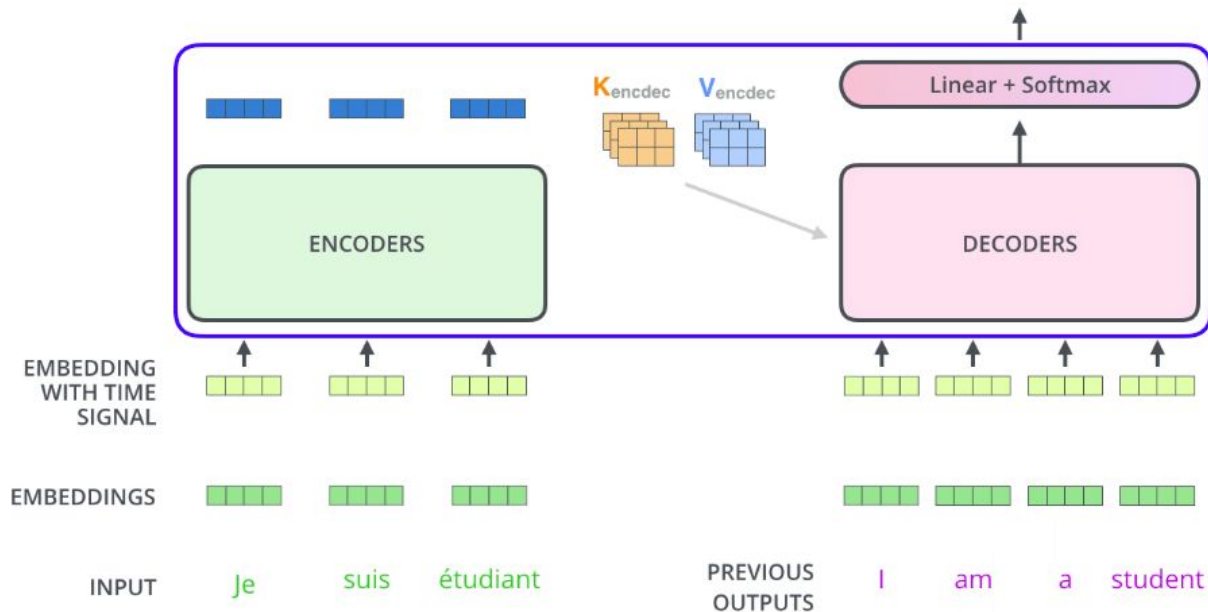
return tf.keras.Model(
    inputs=[inputs, enc_outputs, look_ahead_mask, padding_mask],
    outputs=outputs,
    name=name)
```



Transformer

Decoding time step: 1 2 3 4 5 6

OUTPUT I am a student <end of sentence>





Transformer

```
def transformer(vocab_size,
               num_layers,
               units,
               d_model,
               num_heads,
               dropout,
               name="transformer"):
    inputs = tf.keras.Input(shape=(None,), name="inputs")
    dec_inputs = tf.keras.Input(shape=(None,), name="dec_inputs")

    enc_padding_mask = tf.keras.layers.Lambda(
        create_padding_mask, output_shape=(1, 1, None),
        name='enc_padding_mask')(inputs)
    # Enmascaramos los futuros tokens para la entrada del decoder para el primer bloque de
    look_ahead_mask = tf.keras.layers.Lambda(
        create_look_ahead_mask,
        output_shape=(1, None, None),
        name='look_ahead_mask')(dec_inputs)
    # Enmascamos la salida encoders para el segundo bloque de atención
    dec_padding_mask = tf.keras.layers.Lambda(
        create_padding_mask, output_shape=(1, 1, None),
        name='dec_padding_mask')(inputs)
```



Transformer

```
enc_outputs = encoder(
    vocab_size=vocab_size,
    num_layers=num_layers,
    units=units,
    d_model=d_model,
    num_heads=num_heads,
    dropout=dropout,
)(inputs=[inputs, enc_padding_mask])

dec_outputs = decoder(
    vocab_size=vocab_size,
    num_layers=num_layers,
    units=units,
    d_model=d_model,
    num_heads=num_heads,
    dropout=dropout,
)(inputs=[dec_inputs, enc_outputs, look_ahead_mask, dec_padding_mask])

outputs = tf.keras.layers.Dense(units=vocab_size, name="outputs")(dec_outputs)

return tf.keras.Model(inputs=[inputs, dec_inputs], outputs=outputs, name=name)
```



Entrenamiento



Entrenamiento

Hiper Parámetros

```
NUM_LAYERS = 2
D_MODEL = 256
NUM_HEADS = 8
UNITS = 512
DROPOUT = 0.1
MAX_LENGTH = 40
model = transformer({
    vocab_size=VOCAB_SIZE,
    num_layers=NUM_LAYERS,
    units=UNITS,
    d_model=D_MODEL,
    num_heads=NUM_HEADS,
    dropout=DROPOUT})
```



Entrenamiento

Función de Pérdida

$$CE = - \sum_i^C t_i \log(f(s)_i)$$

```
def loss_function(y_true, y_pred):  
    y_true = tf.reshape(y_true, shape=(-1, MAX_LENGTH - 1))  
  
    loss = tf.keras.losses.SparseCategoricalCrossentropy(  
        from_logits=True, reduction='none')(y_true, y_pred)  
  
    mask = tf.cast(tf.not_equal(y_true, 0), tf.float32)  
    loss = tf.multiply(loss, mask)  
  
    return tf.reduce_mean(loss)
```



Entrenamiento

Diferencia Entre Sparse y Categorical

If your targets are **one-hot encoded**, use `categorical_crossentropy`.

- Examples of **one-hot encodings**:

- `[1,0,0]`
- `[0,1,0]`
- `[0,0,1]`

But if your targets are **integers**, use `sparse_categorical_crossentropy`.

- Examples of integer encodings (*for the sake of completion*):

- `1`
- `2`
- `3`



Entrenamiento

Radio de aprendizaje

```
class CustomSchedule(tf.keras.optimizers.schedules.LearningRateSchedule):  
  
    def __init__(self, d_model, warmup_steps=4000):  
        super(CustomSchedule, self).__init__()  
  
        self.d_model = d_model  
        self.d_model = tf.cast(self.d_model, tf.float32)  
  
        self.warmup_steps = warmup_steps  
  
    def __call__(self, step):  
        arg1 = tf.math.rsqrt(step)  
        arg2 = step * (self.warmup_steps**-1.5)  
  
        return tf.math.rsqrt(self.d_model) * tf.math.minimum(arg1, arg2)
```

Fórmula

$$lrate = d_{\text{model}}^{-0.5} \cdot \min(step_num^{-0.5}, step_num \cdot warmup_steps^{-1.5})$$



Entrenamiento

Radio de aprendizaje


Fórmula

$$lrate = d_{\text{model}}^{-0.5} \cdot \min(step_num^{-0.5}, step_num \cdot warmup_steps^{-1.5})$$

```
In [ ]: sample_learning_rate = CustomSchedule(d_model=128)

plt.plot(sample_learning_rate(tf.range(200000, dtype=tf.float32)))
plt.ylabel("Learning Rate")
plt.xlabel("Train Step")
```

```
Out[ ]: Text(0.5, 0, 'Train Step')
```





Entrenamiento

Optimizador y Métrica

RMSProp

$$\begin{aligned}\theta_t &= \theta_{t-1} - \frac{\eta}{\sqrt{v_t}} g_{t-1} \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2)(g_{t-1})^2 \\ v_1 &= (g_0)^2\end{aligned}$$

SGDM

$$\begin{aligned}\theta_t &= \theta_{t-1} - \eta m_t \\ m_t &= \beta_1 m_{t-1} + (1 - \beta_2) g_{t-1}\end{aligned}$$

Adam

$$\begin{aligned}\theta_t &= \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t}} m_t \text{ (outline)} \\ \theta_t &= \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \text{ (complete)} \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}\end{aligned}$$

```
learning_rate = CustomSchedule(D_MODEL)

optimizer = tf.keras.optimizers.Adam(
    learning_rate, beta_1=0.9, beta_2=0.98, epsilon=1e-9)

def accuracy(y_true, y_pred):
    # ensure labels have shape (batch_size, MAX_LENGTH - 1)
    y_true = tf.reshape(y_true, shape=(-1, MAX_LENGTH - 1))
    return tf.keras.metrics.sparse_categorical_accuracy(y_true, y_pred)
```



Entrenamiento

Optimizador y Métrica

```
def categorical_accuracy(y_true, y_pred):  
    return K.cast(K.equal(K.argmax(y_true, axis=-1),  
                           K.argmax(y_pred, axis=-1)),  
                  K.floatx())  
  
def sparse_categorical_accuracy(y_true, y_pred):  
    return K.cast(K.equal(K.max(y_true, axis=-1),  
                           K.cast(K.argmax(y_pred, axis=-1), K.floatx())),  
                  K.floatx())
```

`categorical_accuracy` checks to see if the *index* of the maximal true value is equal to the *index* of the maximal predicted value.

`sparse_categorical_accuracy` checks to see if the maximal true value is equal to the *index* of the maximal predicted value.

```
learning_rate = CustomSchedule(D_MODEL)  
  
optimizer = tf.keras.optimizers.Adam(  
    learning_rate, beta_1=0.9, beta_2=0.98, epsilon=1e-9)  
  
def accuracy(y_true, y_pred):  
    # ensure labels have shape (batch_size, MAX_LENGTH - 1)  
    y_true = tf.reshape(y_true, shape=(-1, MAX_LENGTH - 1))  
    return tf.keras.metrics.sparse_categorical_accuracy(y_true, y_pred)
```


Revisar documentación de tensorflow “click aquí”



Entrenamiento

```
model.compile(optimizer=optimizer, loss=loss_function, metrics=[accuracy])
```

Python



```
EPOCHS = 15
```

Python

```
model.fit(dataset, epochs=EPOCHS)
```

Python



Entrenamiento

```
Epoch 6/15
690/690 [=====] - 715s 1s/step - loss: 1.2432 - accuracy: 0.0973
Epoch 7/15
690/690 [=====] - 813s 1s/step - loss: 1.1879 - accuracy: 0.1017
Epoch 8/15
690/690 [=====] - 720s 1s/step - loss: 1.1284 - accuracy: 0.1068
Epoch 9/15
690/690 [=====] - 628s 911ms/step - loss: 1.0721 - accuracy: 0.1122
Epoch 10/15
690/690 [=====] - 634s 918ms/step - loss: 1.0196 - accuracy: 0.1180
Epoch 11/15
690/690 [=====] - 617s 894ms/step - loss: 0.9718 - accuracy: 0.1237
Epoch 12/15
690/690 [=====] - 620s 899ms/step - loss: 0.9283 - accuracy: 0.1292
Epoch 13/15

show more (open the raw output data in a text editor) ...

690/690 [=====] - 635s 920ms/step - loss: 0.8881 - accuracy: 0.1344
Epoch 14/15
690/690 [=====] - 631s 915ms/step - loss: 0.8519 - accuracy: 0.1398
Epoch 15/15
690/690 [=====] - 657s 952ms/step - loss: 0.8185 - accuracy: 0.1449

<tensorflow.python.keras.callbacks.History at 0x7f78535f9d10>
```



Evaluación

Función evaluate

```
def evaluate(sentence):
    sentence = preprocess_sentences(sentence)

    sentence = tf.expand_dims(
        START_TOKEN + tokenizer.encode(sentence) + END_TOKEN, axis=0)

    output = tf.expand_dims(START_TOKEN, 0)

    for i in range(MAX_LENGTH):
        predictions = model(inputs=[sentence, output], training=False)
        # Seleccionamos la última palabra de la dimensión seq_len
        predictions = predictions[:, -1:, :]
        predicted_id = tf.cast(tf.argmax(predictions, axis=-1), tf.int32)

        # Retornamos el resultado si predicted_id es igual que el token final.
        if tf.equal(predicted_id, END_TOKEN[0]):
            break

        # Concatenamos el predicted_id a la salida que se le da al decodificador
        # como input.
        output = tf.concat([output, predicted_id], axis=-1)

    return tf.squeeze(output, axis=0)
```



Evaluación

Función predict

```
def predict(sentence, name):  
    prediction = evaluate(sentence)  
  
    predicted_sentence = tokenizer.decode(  
        [i for i in prediction if i < tokenizer.vocab_size])  
  
    print('{}: {}'.format(name, sentence))  
    print('Bot: {}'.format(predicted_sentence))  
  
    return predicted_sentence
```



Resultados

```
output = predict('Hello', 'Alex')
```

Alex: Hello

Bot: hi.



Resultados

```
> output = predict('Where have you been?', "Alex")  
1  
..  
Alex: Where have you been?  
Bot: i m a little scared of myself.
```




Guardar Pesos

```
model.save_weights('path_to_my_model.h5')
```

```
model.load_weights('path_to_my_model.h5')
```



Implementación del Chatbot

```
def dialog(name):  
    while(True):  
        sentence = input('Write a sentences')  
        if sentence.lower() == "f":  
            break  
        predict(sentence, name)
```



Implementación del Chatbot

```
dialog("Alex")
```

Alex: hello

Bot: hi.

Alex: how are you?

Bot: i m sorry i didn t have to do anything. i want to be unzip my dress.

Alex: do you study today?

Bot: i m going to be a little girl.

Alex: tell me about you

Bot: you re not going to be late.

Alex: ok, ok

Bot: you re a fucking asshole! you re a fucking bull to you!

Alex: calm down, please

Bot: i m going to give you something.

Alex: ok, see you

Bot: okay.

Alex: good bye

Bot: thank you.



Conclusiones



Gracias